

Week 3: User Profiles and Content Based RSs

CS3448: Recommender Systems /
CSX4207/ITX4207: Decision Support and
Recommender Systems

Asst. Prof. Dr. Rachsuda Setthawong

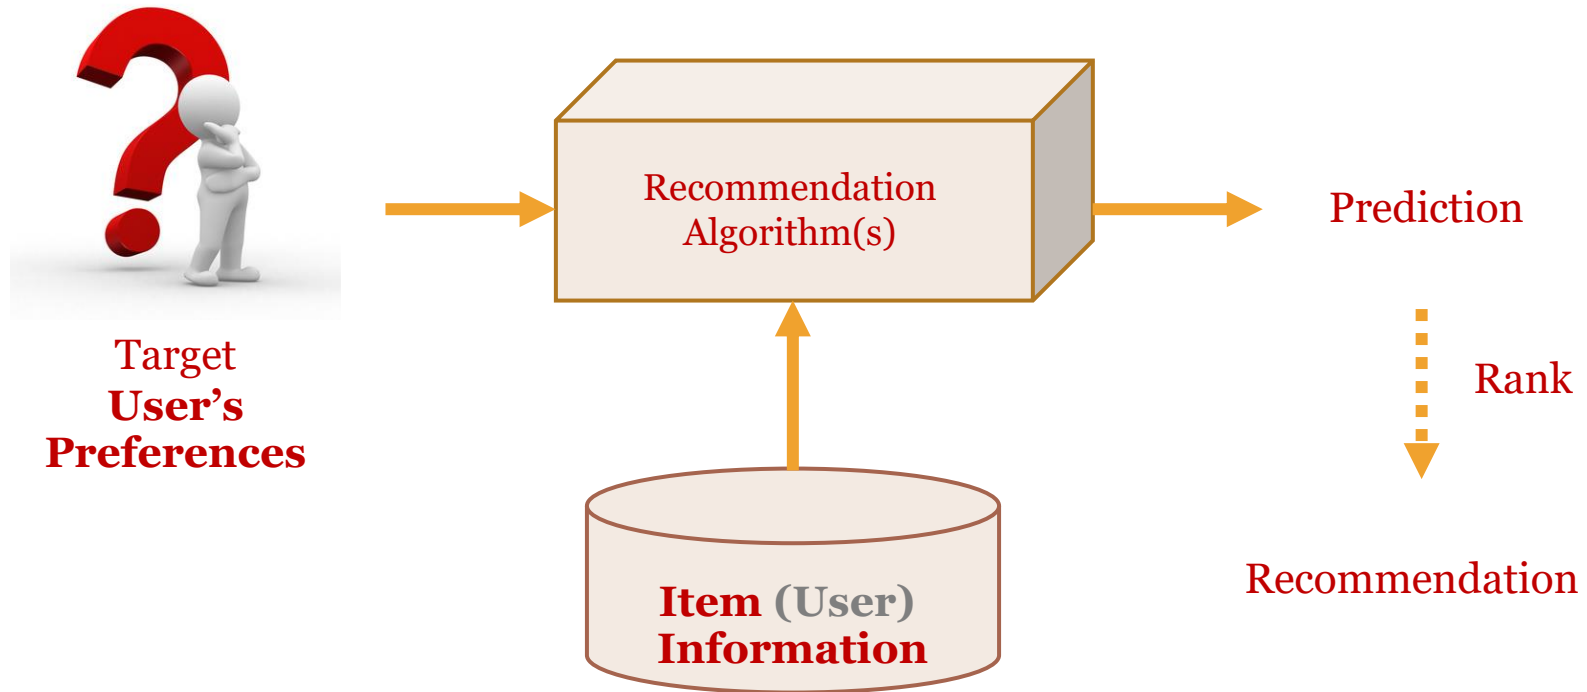
Objectives

- To understand concepts of user profile and be able to construct it
- To understand concepts of content based Recommender Systems and be familiar with some algorithms in this approach
- To understand strong points and weak points of content based RSs

Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case Study
- Available Tools

How to Generate Prediction/Recommendation?



Content-based:

“Display more items similar to what I like.”

Typical User Profile

- The description of what information is of interest to a user
 - An approximation of the real user's interests
- Compact representation in terms of memory and complexity
- Same representation as information filtered

Examples of Items and Descriptors

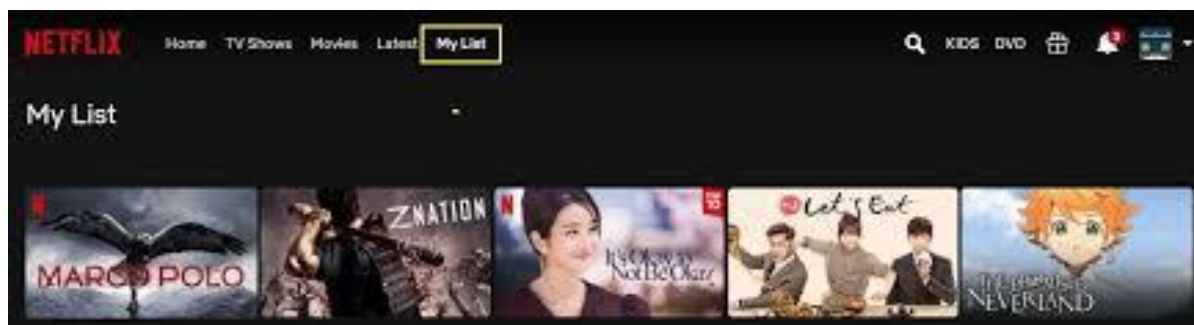
| Item | Descriptors |
|----------------------|--|
| Movie | Genre, Main Actor (Actress), Producer, Production, Release Year |
| Book | Genre, Title, Author(s), Publisher, Abstract, Year |
| Clothes | Fabric Type, Make, Color |
| Restaurant | Type, Rating, Opening Hours, Wifi Availability, Range of Price, Location, Service, Parking Available |
| Vehicle | Type, Make, Model, Color, Horse Power, Number of Doors, Year |
| Music | ?? |
| Attraction (tourism) | ?? |

User Profiling (User Modeling) in RSs

- A process of user profile gathering, construction and representation
- Approaches:
 - **Explicit Model** (ask users straightforwardly)
 - Enabling building and editing profile by user are useful.
 - **Implicit Model** (observe their behaviors, e.g., click, view, buy)
 - May also infer profile from explicit user ratings
 - May require mapping of item preference and attribute preference

A Simple Approach of User Profiling for New Users

- Register for an account and answer couple questions before starting using the websites.
- Prompt users to give ratings on an initial items' set.



User (Item) Representation

- A vector of (features) attributes representing users (items)
 - $v = \langle a_1, a_2, \dots, a_n \rangle$, where v_i , $1 \leq i \leq n$, is an element of $\text{dom}(A_i)$
- Attribute value: $v[A_i]$ or $v \cdot A_i$
 - The i^{th} value in vector v corresponding to attribute A_i .
 - The values can be in form of **weighted** keywords/terms, topics, ratings, etc.
 - Frequency
 - TF-IDF

Example of Vector Construction -- 1

| Vehicle ID | Type | Make | Model | Color | Year |
|------------|--------------|--------|----------|-------|------|
| v_1 | Sedan | Toyota | Altis | White | 2015 |
| v_2 | Sedan | Mazda | 2 | Red | 2014 |
| v_3 | Sedan | Mazda | 2 | White | 2014 |
| v_4 | Wagon | Toyota | Fortuner | Black | 2014 |
| v_5 | Pickup Truck | Toyota | Hilux | Green | 2015 |
| ... | | | | | |

$v_1 = \langle \text{'Sedan'}, \text{'Toyota'}, \text{'Altis'}, \text{'White'}, 2015 \rangle$

$v_2 = \langle \text{'Sedan'}, \text{'Mazda'}, \text{'2'}, \text{'Red'}, 2014 \rangle$

$v_3 = \langle \text{'Sedan'}, \text{'Mazda'}, \text{'2'}, \text{'White'}, 2014 \rangle$

...

Example of Vector Construction -- 2

| Book ID | Title | Genre | Author(s) | Year |
|---------|---|----------|--|------|
| v_1 | The Great Gatsby | Novel | F. Scott Fitzgerald | 1925 |
| v_2 | Lolita | Novel | Vladimir Nabokov | 1955 |
| v_3 | Android Programming: The Big Nerd Ranch Guide | Computer | Brian Hardy, Bill Phillips | 2013 |
| v_4 | Introduction to Data Mining | Computer | Pang-Ning Tan, Michigan State University, Michael Steinbach | 2005 |
| ... | | | | |

$v_1 = \langle \text{'The Great Gatsby'}, \text{'Novel'}, \text{'F. Scott Fitzgerald'}, 1925 \rangle$

$v_2 = \langle \text{'Lolita'}, \text{'Novel'}, \text{'Vladimir Nabokov'}, 1955 \rangle$

$v_3 = \langle \text{'Android Programming: The Big Nerd Ranch Guide'}, \text{'Computer'}, \text{'Brian Hardy, Bill Phillips'}, 2013 \rangle$

...

Note on the Vector of Attributes

- Select only subset attributes that contributes to the recommendation/are applicable to the algorithm used.
- May need to preprocess the data (*e.g.*, *TF-IDF*) to come up with more suitable representation of attribute used.

Cannot be used directly!



| Book ID | Title | Genre | Author(s) | Year |
|---------|------------------|-------|---------------------|------|
| v_1 | The Great Gatsby | Novel | F. Scott Fitzgerald | 1925 |
| v_2 | Lolita | Novel | Vladimir Nabokov | 1955 |
| ... | | | | |

Outlines

- User Profiles and User Profiling
- **Term Frequency and Invert Document Frequency (TF-IDF)**
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case study
- Available Tools

Frequency

- The number of occurrences of a repeating event
(per unit time)

| Book ID | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York |
|---------|------|-------------|--------|-------|-----------|---------|----------|
| V_1 | 50 | 30 | 3 | 12 | 0 | 1 | 15 |
| V_2 | 0 | 0 | 0 | 24 | 0 | 0 | 13 |
| V_3 | 0 | 0 | 25 | 0 | 30 | 55 | 1 |
| V_4 | 0 | 0 | 0 | 0 | 15 | 5 | 0 |

$$v_1 = \langle 50, 30, 3, 12, 0, 1, 15 \rangle$$

$$v_2 = \langle 0, 0, 0, 24, 0, 0, 13 \rangle$$

$$v_3 = \langle 0, 0, 25, 0, 30, 55, 1 \rangle$$

$$v_4 = \langle 0, 0, 0, 0, 15, 5, 0 \rangle$$


Term-Frequency - Inverse Document Frequency (TF-IDF) - 1

- Overcome **the problems of traditional keyword representation** with the following **assumption** *that every word is equally important.*
- Prevent domination of the longer length documents over shorter ones when matching with the user profile.

Term-Frequency - Inverse Document Frequency (TF-IDF) - 2

- TF: Frequency of a term appeared in a document
(Normalization is required.)
- IDF: Decrease the significance of common terms (appeared in all documents.)

TF-IDF



$$TFIDF(t) = TF(t) \cdot IDF(t) \quad (1)$$

A measure calculated from an individual document:

$$TF(t) = \frac{n_{td}}{n_t} \quad (2)$$

A measure calculated from the whole dataset:

$$IDF(t) = \log_e \left(\frac{N_d}{n_{dt}} \right) \quad (3)$$

where,

n_{td} : number of times that a term t appeared in a document d ,

n_t : total number of terms in that document (for normalization),

N_d : total number of documents, and

n_{dt} : number of documents containing the term t .

(if $n_{dt} = 0$, then $IDF(t) = 0$)

consider an individual document

Consider all documents

Asst. Prof. Dr. Rachsuda Setthawong

$$TFIDF(t) = TF(t) \cdot IDF(t) \quad (1)$$

TF-IDF Examples

$$TF(t) = \frac{n_{td}}{n_t} \quad (2)$$

$$IDF(t) = \log_e \left(\frac{N_d}{n_{dt}} \right) \quad (3)$$

Frequency

| Book ID | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|---------|------|-------------|--------|-------|-----------|---------|----------|-----|
| v_1 | 50 | 30 | 3 | 12 | 0 | 1 | 15 | 100 |
| v_2 | 0 | 0 | 0 | 24 | 0 | 0 | 13 | 150 |
| v_3 | 0 | 0 | 25 | 0 | 30 | 55 | 1 | 120 |
| v_4 | 0 | 0 | 0 | 0 | 15 | 5 | 0 | 130 |

TF-IDF

| TFIDF | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|-------|----------|-------------|----------|----------|-----------|----------|----------|-----|
| v_1 | 0.328506 | 0.197103 | 0.009855 | 0.039421 | 0 | 0.001363 | 0.020451 | 0 |
| v_2 | 0 | 0 | 0 | 0.08896 | 0 | 0 | 0.019999 | 0 |
| v_3 | 0 | 0 | 0.075016 | 0 | 0.090019 | 0.068496 | 0.001245 | 0 |
| v_4 | 0 | 0 | 0 | 0 | 0.069315 | 0.009589 | 0 | 0 |

Pros and Cons of TF-IDF

Pros

- Reduce weight of stop words, e.g., 'a', 'an', 'the', 'is', 'am', 'are'
- Increase weight of key terms (not incidental ones)
- Widely used to create a profile of a document
- Feasibly used together with ratings in user profiles

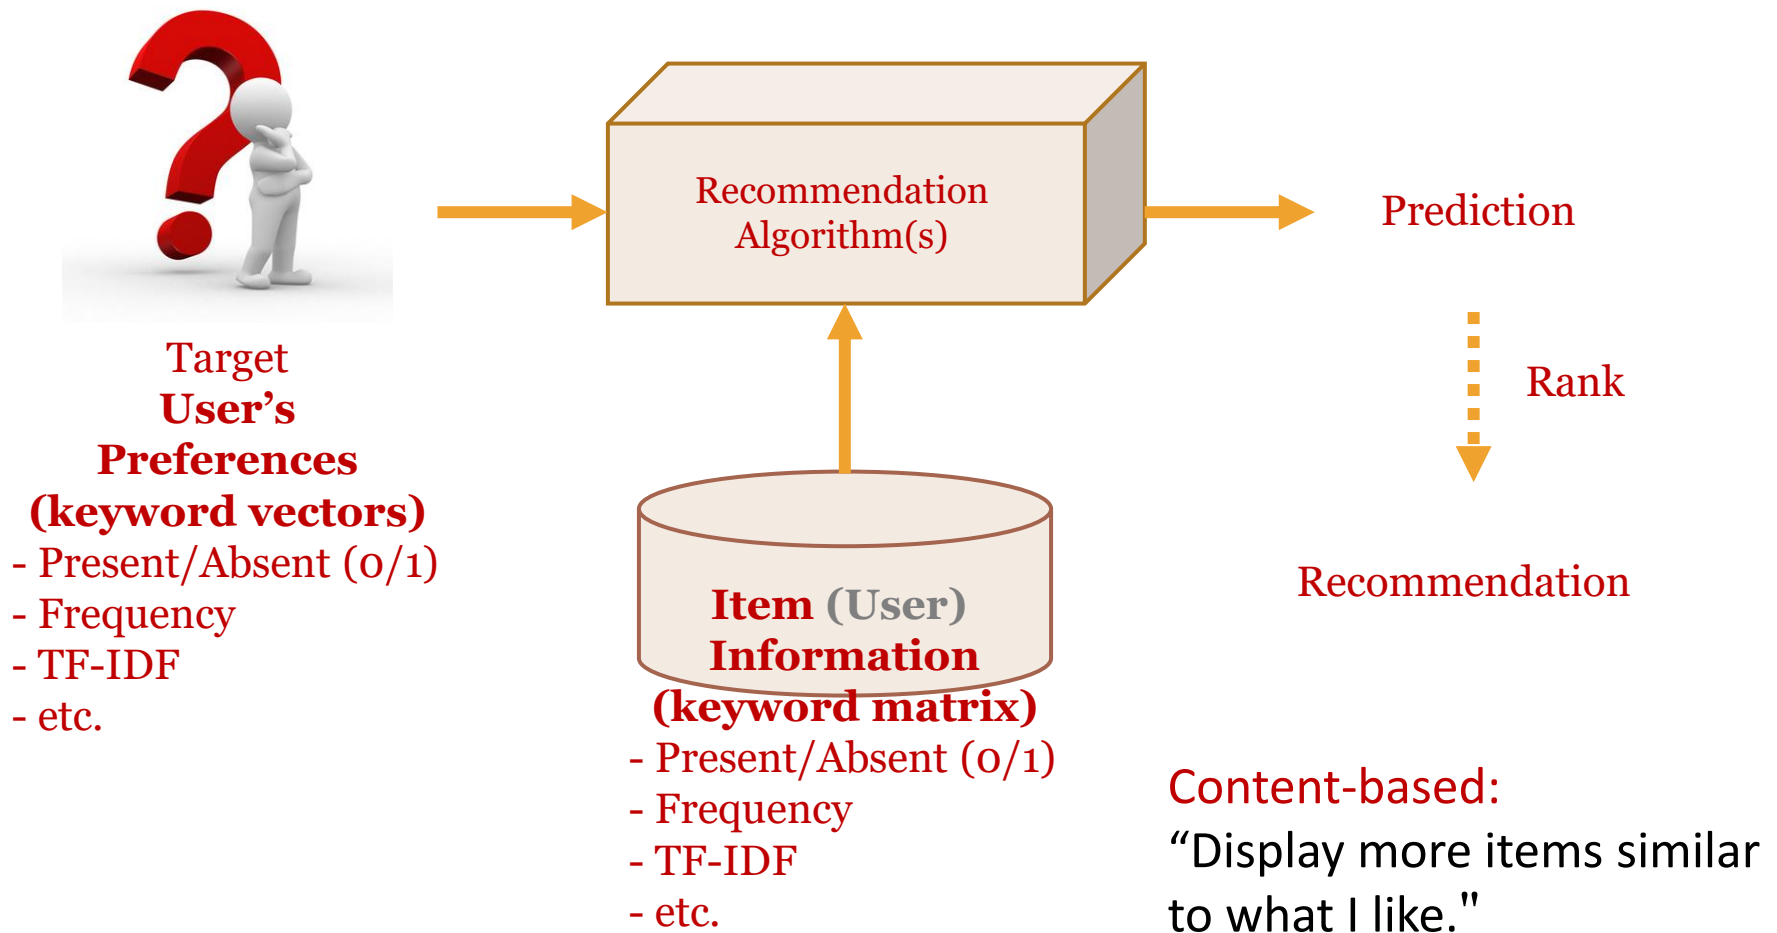
Cons

- Insufficient frequency of key terms could result in low TF-IDF
- Require preprocessing to handle phrases, e.g., 'World War I'
 - N-gram

Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- **How to Generate Recommendation Using Content Based Approach**
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case study
- Available Tools

How to Generate Recommendation Using Content Based Approach - 1



How to Generate Recommendation Using Content Based Approach - 2

1. Constructing user profiles
2. Predicting items
3. Updating user profiles (for system maintenance)

Step 1: Constructing User Profiles - 1

- Observation: a user has experienced several items.
- Commonly used techniques:
 - Aggregation
 - Normalization
 - Weighting

Constructing User Profiles - 2

- A. Simple unary (adding all items)
- B. Unary with threshold (adding items whose rating above threshold)
- C. Weighting with positive rated items (adding all items but with different weights wrt positive rating)
- D. Weighting with positive and negative rated items (adding all items but with different weights wrt positive and negative ratings)

Constructing User Profiles - 3

Example 1

- $p_{\text{Tom}} = \{v_1, v_2\}$

$$v_1 = \langle 50, 30, 3, 12, 0, 1, 15, 100 \rangle \text{ (Frequency)}$$

$$v_2 = \langle 0, 0, 0, 24, 0, 0, 13, 150 \rangle \text{ (Frequency)}$$

- A. Simple unary

$$v_{\text{Tom}} = \langle 50, 30, 3, 36, 0, 1, 28, 250 \rangle$$

Constructing User Profiles - 4

Example 2

- $p_{\text{Tom}} = \{ \langle v_1, 5.0 \rangle, \langle v_2, 2.0 \rangle, \langle v_3, 4.0 \rangle \}$

$$v_1 = \langle 50, 30, 3, 12, 0, 1, 15, 100 \rangle \text{ (Frequency)}$$

$$v_2 = \langle 0, 0, 0, 24, 0, 0, 13, 150 \rangle \text{ (Frequency)}$$

$$v_3 = \langle 20, 0, 2, 14, 0, 0, 10, 120 \rangle \text{ (Frequency)}$$

- B. Unary with threshold (>3.0)

- $v_{\text{Tom}} = \langle 70, 30, 5, 26, 0, 1, 25, 220 \rangle$

Note: Classical Information Retrieval (IR) vs RS

- Classical IR – based methods are based on keywords.
 - Retrieve relevant items based on explicitly searching keywords as user's input
- RS – user profile (preferences) are rather learned than explicitly elicited.
 - Recommend ranked items (without a query) that a target user should like.

Step 2: Predicting Items

- **IDEA:** Calculating similarity between a target user's vector (user profile) and existing items.
- **Similarity measures** for numeric attributes:
 - Distance based similarity
 - e.g., converting Euclidean distance to similarity
 - Cosine similarity
 - Pearson correlation coefficient
 - Etc.

Prediction Example - Step 1

Suppose that $v_{\text{Tom}} = \langle 50, 30, 3, 36, 0, 1, 28, 250 \rangle$ (refer to slide no. 25)

What book to suggest for Tom?

First! Need to calculate TF-IDF for Tom.

TF-IDF_{Tom}

| | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|------------------|----------|-------------|----------|----------|-----------|----------|----------|-----|
| v_{Tom} | 0.174158 | 0.104495 | 0.005225 | 0.062697 | 0 | 0.000723 | 0.020239 | 0 |

TF-IDF of unread book

| TFIDF | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|-------|------|-------------|----------|-------|-----------|----------|----------|-----|
| v_3 | 0 | 0 | 0.075016 | 0 | 0.090019 | 0.068496 | 0.001245 | 0 |
| v_4 | 0 | 0 | 0 | 0 | 0.069315 | 0.009589 | 0 | 0 |

Prediction Example - Step 2

- Distance based similarity (Euclidean)

- Euclidean distance: $distance(v_1, v_2) = \sqrt{\sum_{i \in Item} (v_{1i} - v_{2i})^2}$

- $similarity(v_1, v_2) = \frac{1}{1 + (distance(v_1, v_2))}$

Interpretation:

The closer to 1 the more similar.

- Examples:

- $similarity(v_{Tom}, v_3) = 0.7992$

- $similarity(v_{Tom}, v_4) = 0.8165$**

Suggest the book that more similar to Tom's!

Note: Not suggest the items that are already experienced by the target user




Prediction Example -

Other similarity measure

- Cosine similarity

- A measure of similarity between two non-zero vectors of an inner product space (dot product) that measures the cosine of the angle between them.



- $$\text{cosine_sim}(v_1, v_2) = \frac{v_1 \cdot v_2}{\|v_1\| \|v_2\|}$$

Interpretation:

- Range of possible value is -1 and 1.
- The closer to 1 the more similar.

The dot product of two vectors:

$$v_1 \cdot v_2 = (v_{11} \times v_{21}) + (v_{12} \times v_{22}) + \dots (v_{1n} \times v_{2n})$$

The norm of a vector:

$$\|v\| = \sqrt{a_1^2 + a_2^2 + \dots + a_n^2}$$

Prediction Example -

Other similarity measure (Cont.)


- Examples:
 - **cosine_sim** ($v_{\text{TFIDF_Tom}}$, v_{TFIDF_3}) = **0.0161**
 - cosine_sim ($v_{\text{TFIDF_Tom}}$, v_{TFIDF_4}) = 0.0005

Suggest the book that more similar to Tom's!

Dice Coefficient

(Categorical Attributes)

$$Dice_coefficient(v_1, v_2) = \frac{2 \times |keywords(v_1) \cap keywords(v_2)|}{|keywords(v_1)| + |keywords(v_2)|}$$



Dice Coefficient - Example

Word Occurrence in 2 books

| Book ID | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|---------|------|-------------|--------|-------|-----------|---------|----------|-----|
| v_1 | yes | yes | yes | yes | | yes | yes | yes |
| v_2 | | | | yes | | | yes | yes |

- $$Dice_coefficient(v_1, v_2) = \frac{2 \times |keywords(v_1) \cap keywords(v_2)|}{|keywords(v_1)| + |keywords(v_2)|}$$
- $$= \frac{2 \times 3}{7 + 3} = 0.6$$

Step 3: Updating User Profiles - 1


- Observation: a user buy/use new items from time to time.
- Consequence:
 - Existing user profiles may be outdate.
- Remedy:
 - Allow incremental update of user profiles.
 - Use time decay reflecting item aging to improve accuracy of prediction over time.

Step 3: Updating User Profiles - 2

Time Decay Strategies

- Time Window
 - consider only data in a window (sliding window) that contains either the latest N instances
 - E.g., the latest 1000 ratings
 - consider only the instances contained in the latest time interval
 - E.g., all ratings given in the last 24 hours.

Figure 2: the curves of time function using different T_0



- Time Decay Function
 - Applied as a weight (when calculating similarity (*sim*) between v_1 and v_2 , multiplying *sim* with an exponential decay functions)

$$f(t) = e^{-\lambda \cdot t} \quad \text{where } \lambda = 1/T_0$$

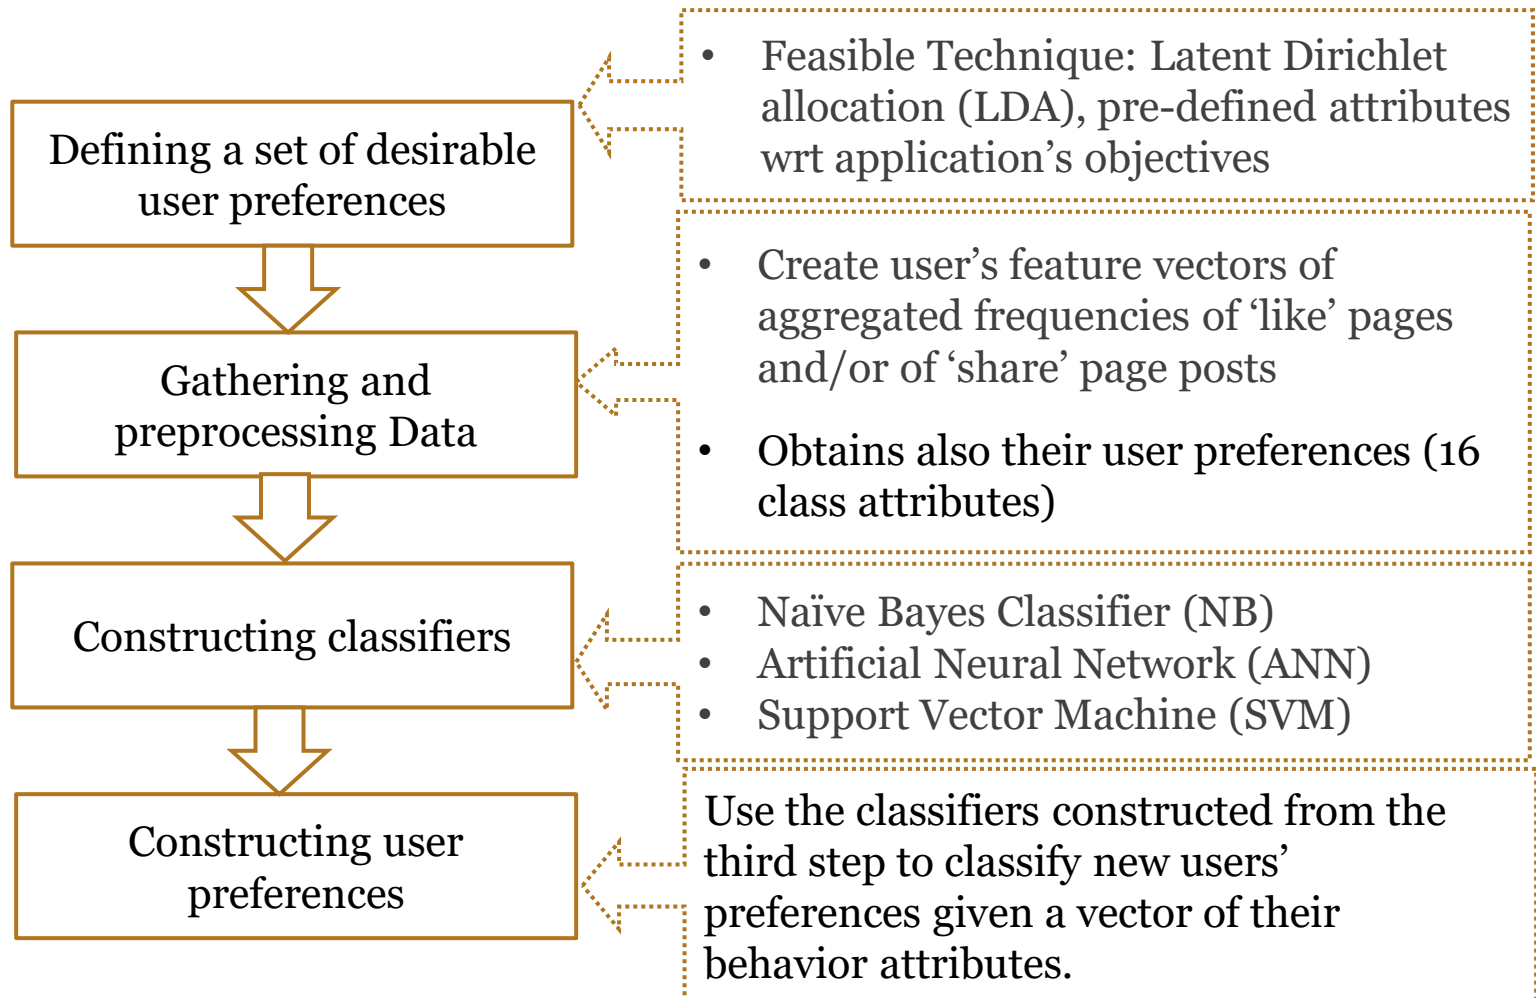
Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case study
- Available Tools


A Technique for User Preference Profiling based on user behaviors on Facebook page categories - 1

- User preferences profiling from social networking data is essential in both social networking mining and recommender systems.
- Facebook provides many useful, both implicit and explicit social data.
- It is a challenge to explore Facebook user behaviors in creating customized user preferences.

A Technique for User Preference Profiling based on User Behaviors on Facebook Page Categories - 2



Percentage of Average Accuracy of 16 User Preferences Using 3 Feature Sets with 3 Algorithms



Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- **Pros and Cons of Content-based RSs**
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case study
- Available Tools

Content Based RSs

Pros

- Easy to compute
- Has structured and available information in some domains, e.g., Book and Movie.
- Easily address new item/user problem (cold start problem)

Cons

- Cannot handle independencies
 - E.g.,
 - I like **Leonardo** **Dicaprio** in 'Catch Me If You Can' but not when he plays 'The Revenant'.
 - I love romantic and comedy movie but not scary, comedy movie.


Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- **Vector Space Model and Recommending Items Using Nearest Neighbors**
- Case study
- Available Tools

The Vector Space Model

- Enhancements
 - Remove stop words , article, preposition, etc ("a", "the", ..)
 - Stemming
 - Use with caution. (stemming of “university” and “universal” is “univers”)
 - User-specified thresholds (e.g., top 100 frequent words or words with TF-IDF greater than its global average)
 - Detection of phrases as terms (such as United Nations)
- Limitations
 - Address neither semantic nor sentiment in texts.
 - Example: My parent do not eat vegetarian since I was very young. So am I.

Recommending Items: Nearest Neighbors - 1



- Given
 - A set of documents D **already rated** by the user (like/dislike)
 - Similarity measure for two document vectors
 - **Input:** vectors of terms in each document,
 - The ratings of each doc can be obtained implicitly or explicitly.
- Find the k nearest neighbors (k-NN) of a not-yet-seen item i in D
 - Take these **ratings** to predict a rating/vote for i
 - Majority voting (If the majority items were liked by the user, so do i .)
 - Variations:
 - neighborhood size
 - lower/upper similarity thresholds
 - Weighting of the votes based on the degree of similarity

An Illustration of Recommending Items: Nearest Neighbors

A Set of Document D

| ItemID | attr1 | attr2 | att3 |
|--------|-------|-------|------|
| 16 | | | |
| 17 | | | |
| 18 | | | |
| ... | | | |
| 100 | | | |

A Target User

| ItemID | attr1 | attr2 | att3 | rating |
|--------|-------|-------|------|--------|
| 1 | | | | 3 |
| 2 | | | | 5 |
| 3 | | | | 4 |
| ... | | | | ... |
| 15 | | | | 5 |

Assumption: rating > 3 implies 'like'
Otherwise, it is 'dislike'

A Similarity Matrix

| ItemID | 16 | 17 | 18 | ... | 100 |
|--------|------------|----|----|-----|-----|
| 1 | sim(16, 1) | | | | |
| 2 | | | | | |
| 3 | | | | | |
| ... | | | | | |
| 15 | | | | | |

Find 3-nn of the item 16, (e.g., the item 1, 2, 15)

Take these ratings to predict a rating/vote for the item 16
= Majority Vote of (2 likes, 1 dislike) → like

An Illustration of Recommending Items: Nearest Neighbors

Variation: neighborhood size

A Set of Document D

| ItemID | attr1 | attr2 | att3 |
|--------|-------|-------|------|
| 16 | | | |
| 17 | | | |
| 18 | | | |
| ... | | | |
| 100 | | | |

A Target User

| ItemID | attr1 | attr2 | att3 | rating |
|--------|-------|-------|------|--------|
| 1 | | | | 3 |
| 2 | | | | 5 |
| 3 | | | | 4 |
| ... | | | | ... |
| 15 | | | | 5 |

Assumption: rating > 3 implies 'like'
Otherwise, it is 'dislike'

A Similarity Matrix

| ItemID | 16 | 17 | 18 | ... | 100 |
|--------|------------|----|----|-----|-----|
| 1 | sim(16, 1) | | | | |
| 2 | | | | | |
| 3 | | | | | |
| ... | | | | | |
| 15 | | | | | |

Find 5-nn of the item 16, (e.g., the item 1, 2, 7, 10, 15)

Take these ratings to predict a rating/vote for the item 16
= Majority Vote of (2 likes, 3 dislikes) → dislike

An Illustration of Recommending Items: Nearest Neighbors

Variation: upper similarity thresholds

(e.g., news domain to select not too similar news)

A Set of Document D

| ItemID | attr1 | attr2 | att3 |
|--------|-------|-------|------|
| 16 | | | |
| 17 | | | |
| 18 | | | |
| ... | | | |
| 100 | | | |

A Target User

| ItemID | attr1 | attr2 | att3 | rating |
|--------|-------|-------|------|--------|
| 1 | | | | 3 |
| 2 | | | | 5 |
| 3 | | | | 4 |
| ... | | | | ... |
| 15 | | | | 5 |

Assumption: rating > 3 implies 'like'
Otherwise, it is 'dislike'

A Similarity Matrix

| ItemID | 16 | 17 | 18 | ... | 100 |
|--------|------------|----|----|-----|-----|
| 1 | sim(16, 1) | | | | |
| 2 | | | | | |
| 3 | | | | | |
| ... | | | | | |
| 15 | | | | | |

Find 3-nn of the item 16, that
Similarity must also be less than 0.9 → (e.g., the item 1, 2, 3)

Sim(16,3) = 0.6, rating 2

Sim(16,1) = 0.6

Sim(16,2) = 0.8

Sim(16,15) = 0.9

Take these ratings to predict a rating/vote for the item 16
= Majority Vote of (1 like, 2 dislikes) → dislike

An Illustration of Recommending Items: Nearest Neighbors

Variation: Weighting of the votes based on the degree of similarity

A Set of Document D

| ItemID | attr1 | attr2 | att3 |
|--------|-------|-------|------|
| 16 | | | |
| 17 | | | |
| 18 | | | |
| ... | | | |
| 100 | | | |

A Target User

| ItemID | attr1 | attr2 | att3 | rating |
|--------|-------|-------|------|--------|
| 1 | | | | 3 |
| 2 | | | | 5 |
| 3 | | | | 4 |
| ... | | | | ... |
| 15 | | | | 5 |

Assumption: rating > 3 implies 'like'
Otherwise, it is 'dislike'

A Similarity Matrix

| ItemID | 16 | 17 | 18 | ... | 100 |
|--------|------------|----|----|-----|-----|
| 1 | sim(16, 1) | | | | |
| 2 | | | | | |
| 3 | | | | | |
| ... | | | | | |
| 15 | | | | | |

Find 3-nn of the item 16, (e.g., the item 1, 2, 15), where

$$\text{Sim}(16, 1) = 0.6$$

$$\text{Sim}(16, 2) = 0.8$$

$$\text{Sim}(16, 15) = 0.9$$

Take these ratings to predict a rating/vote for the item 16
 = **Majority Vote** of (2 likes (0.8+0.9) > 1 dislike (0.6))
 (1.7 > 0.6) → like

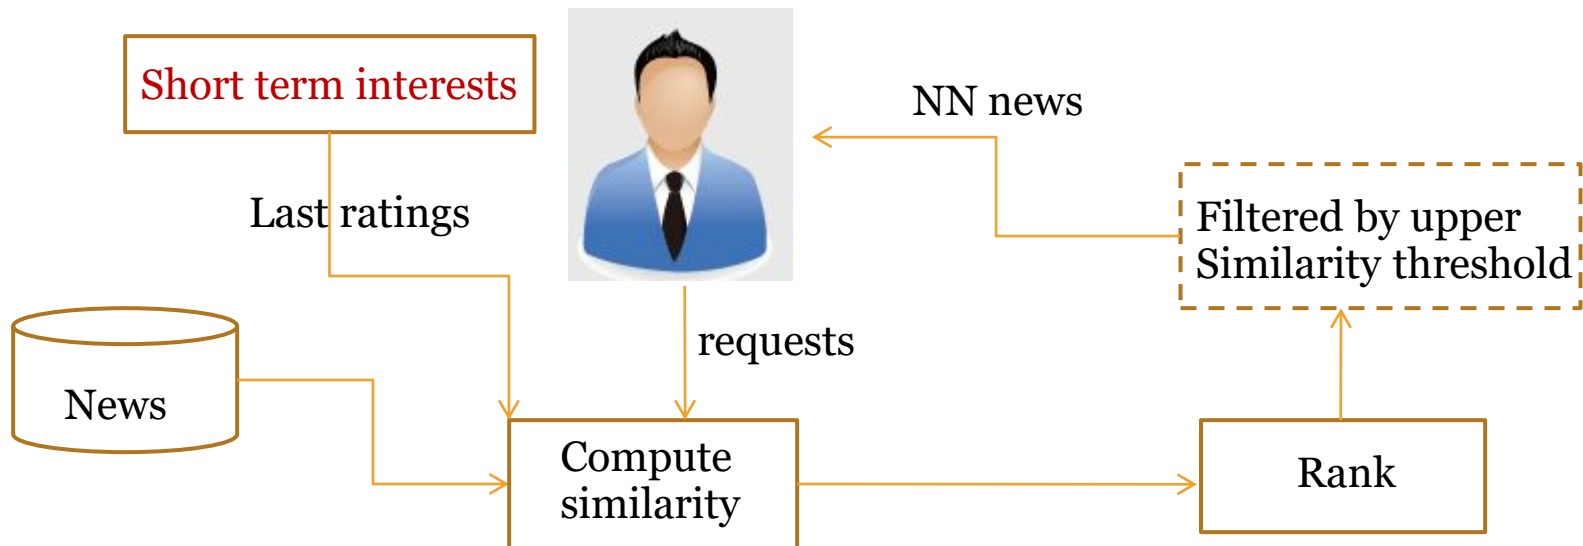
Recommending Items: Nearest Neighbors - 2

- Good to model short-term interests / follow-up stories.
- Used in combination with method to model long-term preferences.

Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- **Case study**
- Available Tools

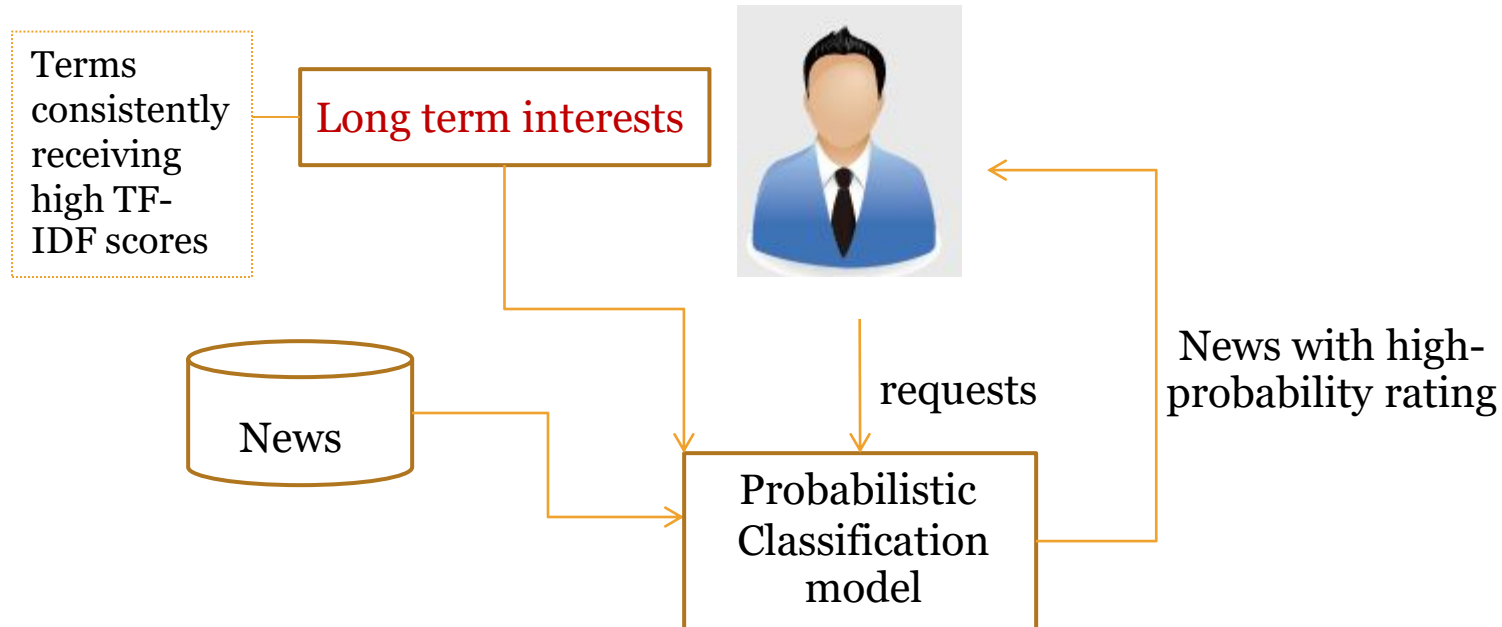
Case Study: the Personalized and Mobile News Access System -- 1 / 3



“Short term interests' model”

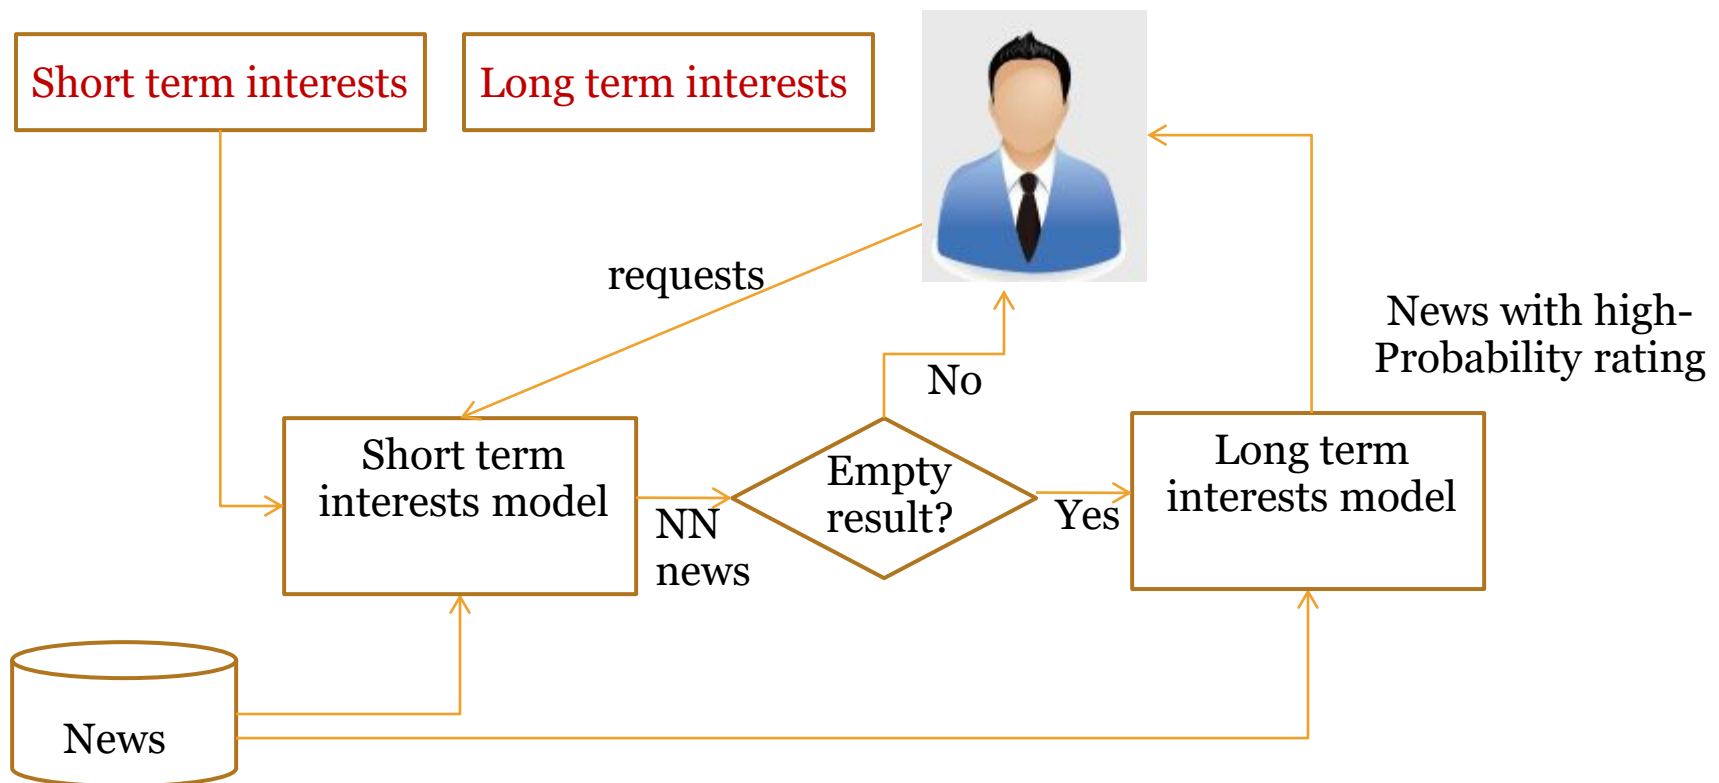
Note: Set an **upper threshold** for *item similarity* to prevent the system from recommending items that the user most probably has already seen.

Case Study: the Personalized and Mobile News Access System -- 2/3



“Long term interests model”

Case Study: the Personalized and Mobile News Access System -- 3/3



“Combined short term and long term interests model”

Advantages of kNN-based Method

- Simple to implement.
- Adapt quickly to recent changes.
- Relatively small number of ratings is sufficient to make a prediction of reasonable quality.

Outlines

- User Profiles and User Profiling
- Term Frequency and Invert Document Frequency (TF-IDF)
- How to Generate Recommendation Using Content Based Approach
 - Additional Similarity Measures
- A Technique for User Preference Profiling based on user behaviors on Facebook page categories
- Pros and Cons of Content-based RSs
- Vector Space Model and Recommending Items Using Nearest Neighbors
- Case study
- Available Tools

Available Tools

- LensKit¹
- Machine Learning Toolkits:
 - Microsoft Machine Learning Studio²
 - Rapidminer Studio³
 - Surprise (scikit) in Python⁴
 - Matlab
 - R
 - Etc.

¹Source: <http://lenskit.org>

²Source: <https://azure.microsoft.com/en-us/blog/building-recommender-systems-with-azure-machine-learning-service/>

³Source: <https://academy.rapidminer.com/learn/article/recommendation-engine-youre-going-to-love-this-movie>

⁴Source: <https://surpriselib.com/>

Practice 3-1

| TFIDF | Yale | World War I | Mining | Hotel | Algorithm | Formula | New York | The |
|-------|----------|-------------|----------|----------|-----------|----------|----------|-----|
| v_1 | 0.328506 | 0.197103 | 0.009855 | 0.039421 | 0 | 0.001363 | 0.020451 | 0 |
| v_2 | 0 | 0 | 0 | 0.08896 | 0 | 0 | 0.019999 | 0 |
| v_3 | 0 | 0 | 0.075016 | 0 | 0.090019 | 0.068496 | 0.001245 | 0 |
| v_4 | 0 | 0 | 0 | 0 | 0.069315 | 0.009589 | 0 | 0 |

- Given $p_{\text{Pete}} = \{v_a, v_b\}$
 - $v_a = \langle 10, 20, 0, 0, 0, 10, 30, 150 \rangle$ (Frequency)
 - $v_b = \langle 0, 5, 20, 5, 0, 0, 10, 100 \rangle$ (Frequency)
- Build user profile vector for Pete (simple unary)
- Predict the book from the example that Pete will like using
 - Cosine similarity
 - Note: use the value of IDF that is already created by the dataset.

Practice 3-2

- Given the dataset below, design and explain user profiles (how feature vectors look like).
- Recommend 3 vehicles to the following user:
 - Mandy is a big fan of Toyota and drove Altis whose model was of year 2010.
- Also explain how you generate the results and show the computation.
 - Similarity measures, item aggregation techniques, etc.

| Vehicle ID | Type | Make | Model | Color | Year |
|------------|--------------|------------|-------------|-------|------|
| v_1 | Sedan | Toyota | Altis | White | 2015 |
| v_2 | Sedan | Mazda | 2 | Red | 2014 |
| v_3 | Sedan | Mazda | 2 | White | 2014 |
| v_4 | Wagon | Toyota | Fortuner | Black | 2014 |
| v_5 | Pickup Truck | Toyota | Hilux | Green | 2015 |
| v_6 | Pickup Truck | Ford | Ranger | Black | 2014 |
| v_7 | Wagon | Mitsubishi | Space Wagon | Black | 2013 |
| v_8 | Sedan | Toyota | Camry | Red | 2015 |
| v_9 | Sedan | Nissan | Almera | White | 2012 |